

Module 3: R Manipulation

Instructor: Anjali Silva, PhD

Data Sciences Institute, University of Toronto

2022

Course Documents

- Visit: <https://github.com/anjalisilva/IntroductionToR>
- All course material will be available via IntroductionToR GitHub repository (<https://github.com/anjalisilva/IntroductionToR>). Folder structure is as follows:
 - Lessons - All files: This folder contains all files.
 - **Lessons - Data only**: This folder contains data only.
 - **Lessons - Lesson Plans only**: This folder contains lesson plans only.
 - **Lessons - PDF only**: This folder contains slide PDFs only.
 - README - README file
 - .gitignore - Files to ignore specified by instructor

Course Contacts

- Instructor: Anjali Silva Email: a.silva@utoronto.ca (Must use the subject line DSI-IntroR. E.g., DSI-IntroR: Inquiry about Lecture I.)
- TA: see GitHub

Overview

- Filtering (Wickham and Grolemund, 2017 Chapter 5, Timbers et al. 2021, Chapter 3.6)
- Arranging (Wickham and Grolemund, 2017 Chapter 5)
- Selecting (Wickham and Grolemund, 2017 Chapter 5, Timbers et al. 2021, Chapter 3.5)
- The pipe (Wickham and Grolemund, 2017 Chapter 5 & 18; Timbers et al. 2021, Chapter 3.8)
- Mutating (Wickham and Grolemund, 2017 Chapter 5, Timbers et al. 2021, Chapter 3.7, 3.10)
- Summarising (Wickham and Grolemund, 2017 Chapter 5, Timbers et al. 2021, Chapter 3.9)
- Grouping (Wickham and Grolemund, 2017 Chapter 5)
- Cleaning (Alexander, 2022, Chapter 11)

Take a look

```
glimpse(ads_data)
```

```
## Rows: 1,460
## Columns: 52
## $ StartDate      <dtm> 2019-06-14 09:43:20,...
## $ EndDate        <dtm> 2019-06-14 09:44:30,...
## $ Status         <dbl+lbl> 0, 0, 0, 0, 0, 0,...
## $ Progress       <dbl> 100, 100, 100, 100, 1...
## $ Duration__in_seconds_ <dbl> 70, 105, 88, 109, 109...
## $ Finished       <dbl+lbl> 1, 1, 1, 1, 1, 1,...
## $ RecordedDate   <dtm> 2019-06-14 09:44:31,...
## $ ResponseId     <chr> "R_11dq3s9btLX57LD", ...
## $ DistributionChannel <chr> "anonymous", "anonymo...
## $ UserLanguage   <chr> "EN", "EN", "EN", "EN...
## $ Consent        <dbl+lbl> 1, 1, 1, 1, 1, 1,...
## $ Pol_7          <dbl+lbl> 5, 3, 1, 2, 6, 4,...
## $ W2_Knowledge   <dbl+lbl> 2, 2, 4, 1, 3, 2,...
## $ Gender         <dbl+lbl> 2, 1, 2, 1, 1, 1,...
## $ Race           <dbl+lbl> 1, 1, 1, 1, 1, 3,...
## $ W1_Feeling_1   <dbl> 2, 1, 4, 3, 3, 3, 6, ...
## $ W1_Actions_1_1 <dbl+lbl> NA, NA, NA, NA, N...
## $ W1_Actions_1_2 <dbl+lbl> 1, NA, NA, 1, N...
## $ W1_Actions_1_3 <dbl+lbl> NA, NA, 1, NA, N...
```

Filtering

Filtering allows us to select rows based on specific traits

```
filter(ads_data, Duration__in_seconds_ < 100)
```

```
## # A tibble: 41 × 52
##   StartDate      EndDate      Status
##   <dtm>         <dtm>         <dbl+lbl>
## 1 2019-06-14 09:43:20 2019-06-14 09:44:30 0 [IP Add...
## 2 2019-06-14 09:43:29 2019-06-14 09:44:58 0 [IP Add...
## 3 2019-06-14 09:44:00 2019-06-14 09:45:11 0 [IP Add...
## 4 2019-06-14 09:43:32 2019-06-14 09:45:12 0 [IP Add...
## 5 2019-06-14 09:43:48 2019-06-14 09:45:25 0 [IP Add...
## 6 2019-06-14 09:44:24 2019-06-14 09:45:26 0 [IP Add...
## 7 2019-06-14 09:43:50 2019-06-14 09:45:29 0 [IP Add...
## 8 2019-06-14 09:44:15 2019-06-14 09:45:42 0 [IP Add...
## 9 2019-06-14 09:44:30 2019-06-14 09:45:58 0 [IP Add...
## 10 2019-06-14 09:44:36 2019-06-14 09:46:05 0 [IP Add...
## # ... with 31 more rows, and 49 more variables:
## #   Progress <dbl>, Duration__in_seconds_ <dbl>,
## #   Finished <dbl+lbl>, RecordedDate <dtm>,
## #   ResponseId <chr>, DistributionChannel <chr>,
## #   UserLanguage <chr>, Consent <dbl+lbl>,
## #   Pol_7 <dbl+lbl>, W2_Knowledge <dbl+lbl>,
```

Arranging

Arranging allows us to sort the order of the table by a certain column

```
arrange(ads_data, Duration__in_seconds_)
```

```
## # A tibble: 1,460 × 52
##   StartDate      EndDate      Status
##   <dtm>         <dtm>         <dbl+lbl>
## 1 2019-06-14 09:58:11 2019-06-14 09:59:01 0 [IP Add...
## 2 2019-06-14 09:44:24 2019-06-14 09:45:26 0 [IP Add...
## 3 2019-06-14 09:43:20 2019-06-14 09:44:30 0 [IP Add...
## 4 2019-06-14 09:44:00 2019-06-14 09:45:11 0 [IP Add...
## 5 2019-06-14 09:52:10 2019-06-14 09:53:26 0 [IP Add...
## 6 2019-06-14 09:45:57 2019-06-14 09:47:13 0 [IP Add...
## 7 2019-06-14 09:50:37 2019-06-14 09:51:53 0 [IP Add...
## 8 2019-06-14 09:45:49 2019-06-14 09:47:08 0 [IP Add...
## 9 2019-06-14 10:10:25 2019-06-14 10:11:45 0 [IP Add...
## 10 2019-06-14 09:53:33 2019-06-14 09:54:54 0 [IP Add...
## # ... with 1,450 more rows, and 49 more variables:
## #   Progress <dbl>, Duration__in_seconds_ <dbl>,
## #   Finished <dbl+lbl>, RecordedDate <dtm>,
## #   ResponseId <chr>, DistributionChannel <chr>,
## #   UserLanguage <chr>, Consent <dbl+lbl>,
## #   Pol_7 <dbl+lbl>, W2_Knowledge <dbl+lbl>,
```

Selecting

Selecting allows us to pick certain columns

```
select(ads_data, RecordedDate)
```

```
## # A tibble: 1,460 × 1
##   RecordedDate
##   <dtm>
## 1 2019-06-14 09:44:31
## 2 2019-06-14 09:44:58
## 3 2019-06-14 09:44:59
## 4 2019-06-14 09:45:00
## 5 2019-06-14 09:45:01
## 6 2019-06-14 09:45:12
## 7 2019-06-14 09:45:12
## 8 2019-06-14 09:45:13
## 9 2019-06-14 09:45:13
## 10 2019-06-14 09:45:16
## # ... with 1,450 more rows
```

Selecting

We can also remove columns

```
select(ads_data, -Consent, -DistributionChannel)
```

```
## # A tibble: 1,460 × 50
##   StartDate      EndDate      Status
##   <dtm>         <dtm>         <dbl+lbl>
## 1 2019-06-14 09:43:20 2019-06-14 09:44:30 0 [IP Add...
## 2 2019-06-14 09:43:11 2019-06-14 09:44:57 0 [IP Add...
## 3 2019-06-14 09:43:29 2019-06-14 09:44:58 0 [IP Add...
## 4 2019-06-14 09:43:10 2019-06-14 09:45:00 0 [IP Add...
## 5 2019-06-14 09:43:11 2019-06-14 09:45:00 0 [IP Add...
## 6 2019-06-14 09:44:00 2019-06-14 09:45:11 0 [IP Add...
## 7 2019-06-14 09:43:32 2019-06-14 09:45:12 0 [IP Add...
## 8 2019-06-14 09:43:27 2019-06-14 09:45:12 0 [IP Add...
## 9 2019-06-14 09:43:08 2019-06-14 09:45:13 0 [IP Add...
## 10 2019-06-14 09:43:36 2019-06-14 09:45:16 0 [IP Add...
## # ... with 1,450 more rows, and 47 more variables:
## #   Progress <dbl>, Duration__in_seconds_ <dbl>,
## #   Finished <dbl+lbl>, RecordedDate <dtm>,
## #   ResponseId <chr>, UserLanguage <chr>,
## #   Pol_7 <dbl+lbl>, W2_Knowledge <dbl+lbl>,
## #   Gender <dbl+lbl>, Race <dbl+lbl>,
```


The pipe

So far, we have written our code like this:

```
filter(ads_data, Duration__in_seconds_ < 100)
```

```
## # A tibble: 41 × 52
##   StartDate      EndDate      Status
##   <dtm>         <dtm>         <dbl+lbl>
## 1 2019-06-14 09:43:20 2019-06-14 09:44:30 0 [IP Add...
## 2 2019-06-14 09:43:29 2019-06-14 09:44:58 0 [IP Add...
## 3 2019-06-14 09:44:00 2019-06-14 09:45:11 0 [IP Add...
## 4 2019-06-14 09:43:32 2019-06-14 09:45:12 0 [IP Add...
## 5 2019-06-14 09:43:48 2019-06-14 09:45:25 0 [IP Add...
## 6 2019-06-14 09:44:24 2019-06-14 09:45:26 0 [IP Add...
## 7 2019-06-14 09:43:50 2019-06-14 09:45:29 0 [IP Add...
## 8 2019-06-14 09:44:15 2019-06-14 09:45:42 0 [IP Add...
## 9 2019-06-14 09:44:30 2019-06-14 09:45:58 0 [IP Add...
## 10 2019-06-14 09:44:36 2019-06-14 09:46:05 0 [IP Add...
## # ... with 31 more rows, and 49 more variables:
## #   Progress <dbl>, Duration__in_seconds_ <dbl>,
## #   Finished <dbl+lbl>, RecordedDate <dtm>,
## #   ResponseId <chr>, DistributionChannel <chr>,
## #   UserLanguage <chr>, Consent <dbl+lbl>,
## #   Pol_7 <dbl+lbl>, W2_Knowledge <dbl+lbl>,
```

The pipe

We can use the pipe `%>%`, which passes what we wrote on the previous line into the next function as the first argument:

```
ads_data %>%  
  filter(Duration__in_seconds_ < 100) %>%  
  arrange(Duration__in_seconds_) %>%  
  select(RecordedDate, Duration__in_seconds_)
```

```
## # A tibble: 41 × 2  
##   RecordedDate      Duration__in_seconds_  
##   <dtm>              <dbl>  
## 1 2019-06-14 09:59:02          50  
## 2 2019-06-14 09:45:26          61  
## 3 2019-06-14 09:44:31          70  
## 4 2019-06-14 09:45:12          70  
## 5 2019-06-14 09:53:26          75  
## 6 2019-06-14 09:47:13          76  
## 7 2019-06-14 09:51:54          76  
## 8 2019-06-14 09:47:08          78  
## 9 2019-06-14 10:11:46          79  
## 10 2019-06-14 09:54:54          80  
## # ... with 31 more rows
```

The pipe

```
ads_data %>%  
  filter(Duration__in_seconds_ < 100) %>%  
  arrange(Duration__in_seconds_) %>%  
  select(RecordedDate, Duration__in_seconds_)
```

You can think of this like:

- Take the ADS data
- Filter so we only have the rows where the survey duration is less than 100 seconds
- Arrange so we go from lowest duration to highest
- Select only the date recorded and the duration

Mutating

Mutating can be used to create new columns or change existing columns.

```
ads_data <- ads_data %>%  
  mutate(Birthyear_add_day = str_c(Birthyear, "07-01")) %>%  
  mutate(Birthyear_add_day = as_datetime(Birthyear_add_day))
```

```
## # A tibble: 1,460 × 3  
##      EndDate          Birthyear Birthyear_add_day  
##      <dtm>          <dbl> <dtm>  
## 1 2019-06-14 09:44:30      1993 1993-07-01 00:00:00  
## 2 2019-06-14 09:44:57      1978 1978-07-01 00:00:00  
## 3 2019-06-14 09:44:58      1993 1993-07-01 00:00:00  
## 4 2019-06-14 09:45:00      1983 1983-07-01 00:00:00  
## 5 2019-06-14 09:45:00      1990 1990-07-01 00:00:00  
## 6 2019-06-14 09:45:11      1980 1980-07-01 00:00:00  
## 7 2019-06-14 09:45:12      1996 1996-07-01 00:00:00  
## 8 2019-06-14 09:45:12      1986 1986-07-01 00:00:00  
## 9 2019-06-14 09:45:13      2000 2000-07-01 00:00:00  
## 10 2019-06-14 09:45:16      1988 1988-07-01 00:00:00  
## # ... with 1,450 more rows
```

Mutating

```
ads_data %>%  
  mutate(age = EndDate - Birthyear_add_day)
```

```
## # A tibble: 1,460 × 4  
##   EndDate      Birthyear Birthyear_add_day  
##   <dtm>      <dbl> <dtm>  
## 1 2019-06-14 09:44:30      1993 1993-07-01 00:00:00  
## 2 2019-06-14 09:44:57      1978 1978-07-01 00:00:00  
## 3 2019-06-14 09:44:58      1993 1993-07-01 00:00:00  
## 4 2019-06-14 09:45:00      1983 1983-07-01 00:00:00  
## 5 2019-06-14 09:45:00      1990 1990-07-01 00:00:00  
## 6 2019-06-14 09:45:11      1980 1980-07-01 00:00:00  
## 7 2019-06-14 09:45:12      1996 1996-07-01 00:00:00  
## 8 2019-06-14 09:45:12      1986 1986-07-01 00:00:00  
## 9 2019-06-14 09:45:13      2000 2000-07-01 00:00:00  
## 10 2019-06-14 09:45:16      1988 1988-07-01 00:00:00  
## # ... with 1,450 more rows, and 1 more variable:  
## #   age <drtn>
```

Summary

```
summary(ads_data)
```

```
##      StartDate
##  Min.      :2019-06-14 09:43:03.00
## 1st Qu.:2019-06-14 09:46:47.50
## Median :2019-06-14 09:52:50.00
## Mean    :2019-06-14 09:57:40.11
## 3rd Qu.:2019-06-14 10:06:28.25
## Max.    :2019-06-14 11:19:45.00
##
##      EndDate      Status
##  Min.      :2019-06-14 09:44:30.00  Min.      :0
## 1st Qu.:2019-06-14 09:51:29.00  1st Qu.:0
## Median :2019-06-14 09:57:57.00  Median :0
## Mean    :2019-06-14 10:02:23.89  Mean     :0
## 3rd Qu.:2019-06-14 10:11:19.50  3rd Qu.:0
## Max.    :2019-06-14 11:27:10.00  Max.     :0
##
##      Progress  Duration__in_seconds_  Finished
##  Min.      :100  Min.      : 50.0  Min.      :1
## 1st Qu.:100  1st Qu.: 178.0  1st Qu.:1
## Median :100  Median : 237.0  Median :1
## Mean     :100  Mean     : 283.3  Mean     :1
```

Pulling a variable for calculations

```
ads_data %>%  
  pull(Duration__in_seconds_)
```

```
##      [1]      70      105      88      109      109      70      99      105      124  
##     [10]     100      96     102      61      98     120      86     119     120  
##     [19]     143     115     131     164     140     126      88     127     146  
##     [28]      88     134     163     111     164     123     176     102     119  
##     [37]     187     179     140     144     183     139     123     162     152  
##     [46]     184     160     181     163     168     101     190     178     144  
##     [55]     194     123     133     135     185     121     163     192     210  
##     [64]     167     139     204     117     170     170     199      95     126  
##     [73]     208     178     207     146     118     170     110     172     226  
##     [82]      78     160     185     186     222     212     185     168     213  
##     [91]      76     213     165     173     218     207     214     203     206  
##    [100]     213     228     186     240     248     208     176     217     142  
##    [109]     190     215     247     163     239     251     185     176     217  
##    [118]     193     171     159     239     252     178     168     101     213  
##    [127]     227     122     217     225     239     182     178     165     248  
##    [136]     190     272     222     101     173     270     121     191     275  
##    [145]     210     227     283     188     194     275     236     169     151  
##    [154]     295     262     257     234     119     287     276     264     286
```

Using the pulled variable for descriptive statistics

Median

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  median(na.rm = TRUE)
```

```
## [1] 237
```

We have to tell the `mean()` function to disregard NAs by writing `na.rm = TRUE`

Using the pulled variable for descriptive statistics

Mean

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  mean(na.rm = TRUE)
```

```
## [1] 283.261
```

Using the pulled variable for descriptive statistics

Range can be calculated using the `range()` function.

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  range(na.rm = TRUE)
```

```
## [1]    50 1575
```

Variance can be calculated using the `var()` function.

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  var(na.rm = TRUE)
```

```
## [1] 29487.81
```

Using the pulled variable for descriptive statistics

Standard Deviation can be calculated using the `sd()` function.

```
ads_data %>%  
  pull(Duration__in_seconds_) %>%  
  sd(na.rm = TRUE)
```

```
## [1] 171.7202
```

Summarise

```
ads_data %>%  
  summarise(mean_time = mean(Duration__in_seconds_, na.rm = TRUE),  
            sd_time = sd(Duration__in_seconds_, na.rm = TRUE))
```

```
## # A tibble: 1 × 2  
##   mean_time sd_time  
##   <dbl>    <dbl>  
## 1     283.     172.
```

Grouping

Before summarising, we can group by a categorical variable

```
ads_data %>%  
  group_by(Gender) %>%  
  summarise(count = n(),  
            mean_time = mean(Duration__in_seconds_, na.rm = TRUE),  
            sd_time = sd(Duration__in_seconds_, na.rm = TRUE))
```

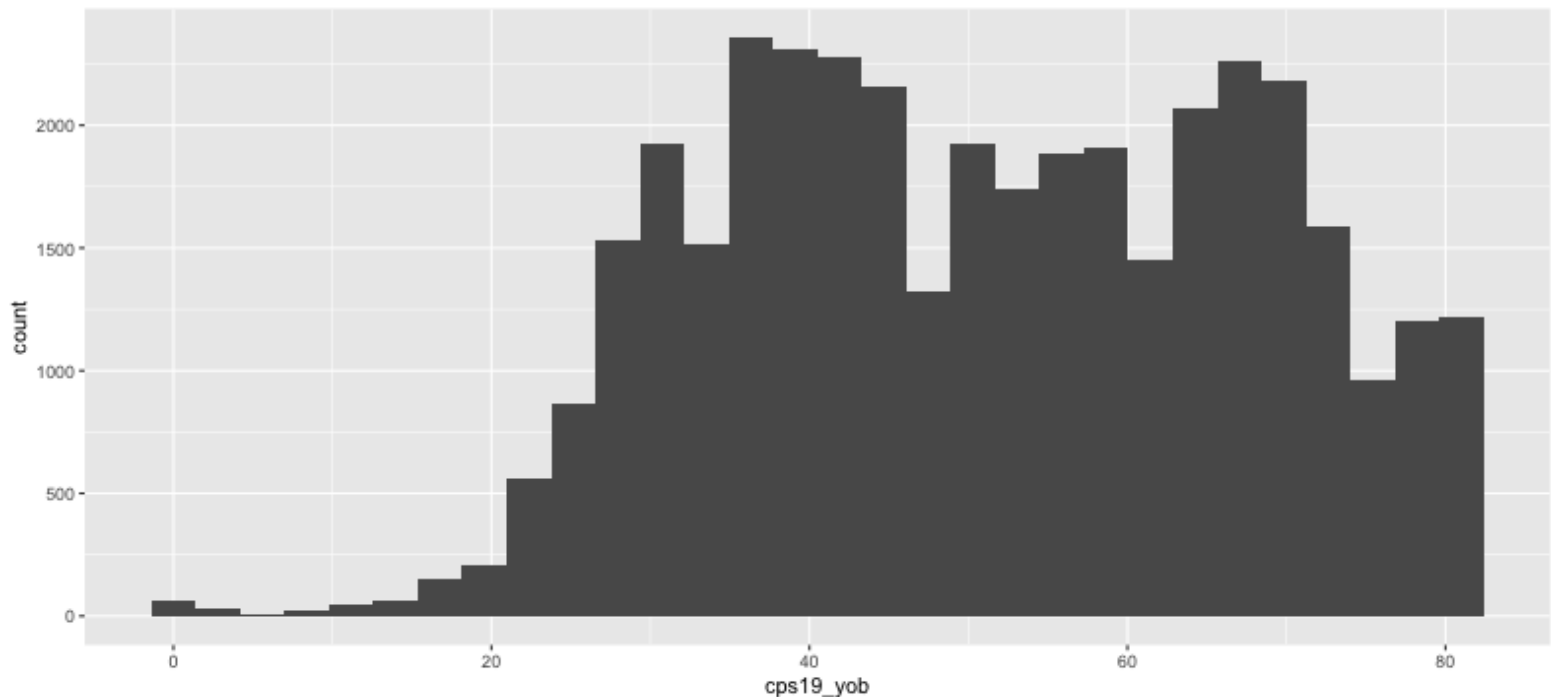
```
## # A tibble: 3 × 4  
##   Gender          count mean_...1 sd_time  
##   <dbl+lbl>      <int>    <dbl>    <dbl>  
## 1 1 [Male]         758      269.    162.  
## 2 2 [Female]       698      299.    181.  
## 3 3 [Prefer a third option/Oth...    4      229     37.7  
## # ... with abbreviated variable name 'mean_time'
```

Manipulation application: data cleaning

Data cleaning

Graphing year of birth shows that it goes from 1 to about 80.

```
ces_2019_raw %>%  
  ggplot(aes(x = cps19_yob)) +  
  geom_histogram()
```



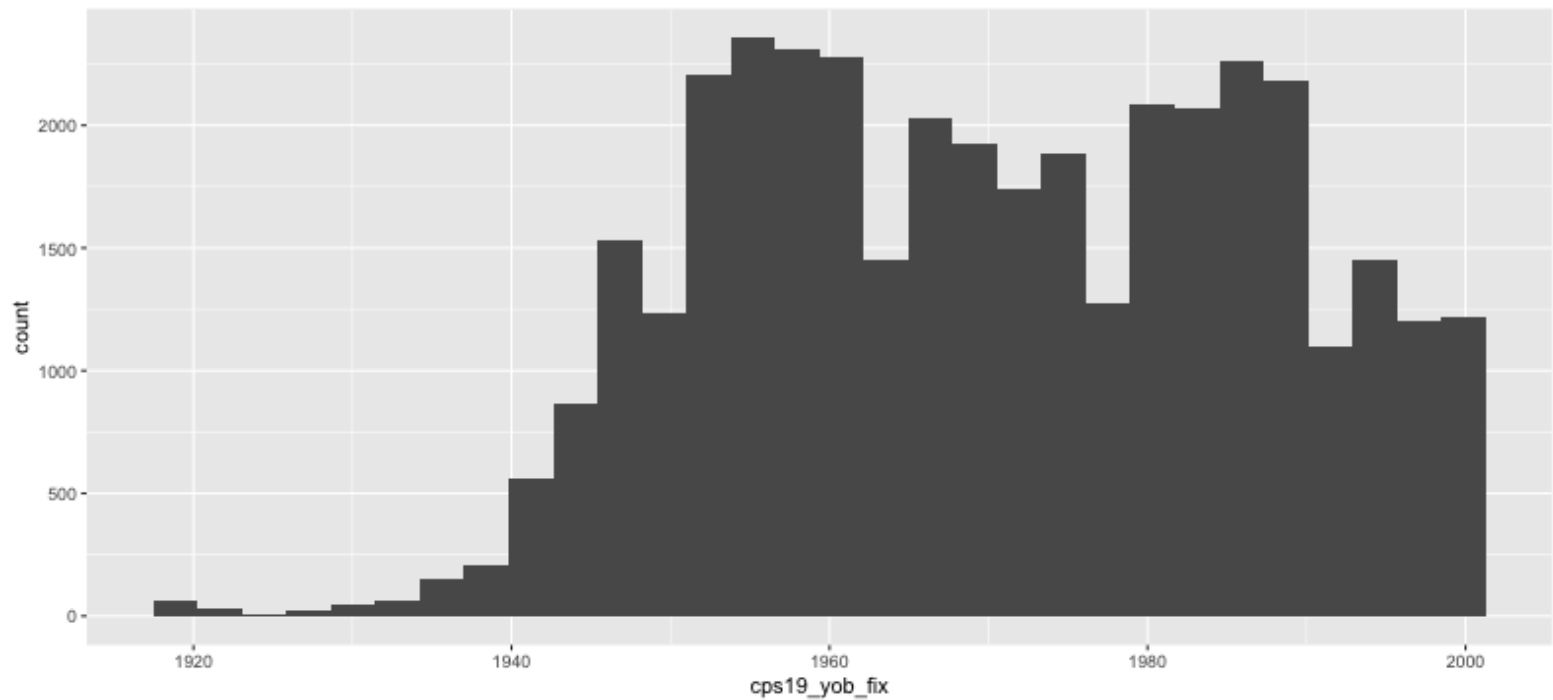
Data cleaning

The codebook says that a value of 1 corresponds to a birth year of 1920, value of 2 to a birth year of 1921, and so on. We can create a new variable that reads more intuitively.

```
CES_data <- ces_2019_raw %>%  
  mutate(cps19_yob_fix = cps19_yob + 1919)
```


Data cleaning

```
CES_data %>%  
  ggplot(aes(x = cps19_yob_fix)) +  
  geom_histogram()
```



Better!


Add a variable for age

Now that we have an accurate birth year, maybe we would like to have the age of the individual as well.

```
CES_data <- CES_data %>%  
  mutate(age = 2019 - cps19_yob_fix)
```


Add a variable for age

```
CES_data %>%  
  ggplot(aes(x = age)) +  
  geom_histogram()
```



Recoding the gender variable

```
CES_data %>%  
  ggplot(aes(x = cps19_gender)) +  
  geom_bar()
```




Recoding the gender variable

```
CES_data <- CES_data %>%  
  mutate(cps19_gender_fix = factor(cps19_gender)) %>%  
  mutate(cps19_gender_fix = fct_recode(cps19_gender_fix,  
                                       "M" = "1",  
                                       "F" = "2",  
                                       "NB" = "3"))
```

Recoding the gender variable

```
CES_data %>%  
  ggplot(aes(x = cps19_gender_fix)) +  
  geom_bar()
```



Fixing household counts

```
CES_data %>%  
  filter(cps19_household > 10) %>%  
  arrange(-cps19_household) %>%  
  pull(cps19_household)
```

```
## [1] 7766666 72000 50000 20000 10000 5667  
## [7] 2000 501 321 99 89 87  
## [13] 69 54 54 50 44 40  
## [19] 34 33 29 27 23 22  
## [25] 22 20 20 20 15 15  
## [31] 13 13 12 12 12 11  
## [37] 11 11 11 11 11 11  
## [43] 11
```

Fixing household counts

```
CES_data <- CES_data %>%  
  mutate(cps19_household = ifelse(cps19_household > 15,  
                                   NA,  
                                   cps19_household))
```

```
CES_data %>%  
  filter(cps19_household > 10) %>%  
  pull(cps19_household)
```

```
## [1] 12 11 15 12 11 13 11 11 11 15 13 12 11 11 11
```


Fixing income

```
CES_data %>%  
  filter(cps19_income_number > 1000000) %>%  
  arrange(-cps19_income_number) %>%  
  pull(cps19_income_number)
```

```
## [1] 6.747658e+60 1.000000e+21 1.000000e+15  
## [4] 8.769655e+10 8.889899e+09 3.062936e+09  
## [7] 1.000000e+09 1.000000e+09 6.788765e+08  
## [10] 3.000000e+08 7.245600e+07 3.454534e+07  
## [13] 3.000000e+07 1.000000e+07 9.999999e+06  
## [16] 8.900000e+06 7.696588e+06 7.440000e+06  
## [19] 6.848382e+06 6.787145e+06 6.782800e+06  
## [22] 6.500100e+06 4.500000e+06 3.000000e+06  
## [25] 2.332100e+06 2.000000e+06 2.000000e+06  
## [28] 1.872717e+06 1.800000e+06 1.650000e+06  
## [31] 1.500000e+06 1.500000e+06 1.450000e+06  
## [34] 1.300000e+06 1.290000e+06 1.250000e+06  
## [37] 1.250000e+06 1.250000e+06 1.150000e+06
```

Fixing income

```
CES_data <- CES_data %>%  
  mutate(cps19_income_number = ifelse(cps19_income_number >= 1000000000,  
                                       NA,  
                                       cps19_income_number))  
  
CES_data %>%  
  filter(cps19_income_number > 1000000) %>%  
  pull(cps19_income_number)
```

```
## [1] 2000000 1500000 4500000 3000000  
## [5] 6848382 7696588 6787145 1250000  
## [9] 1650000 1872717 678876545 1300000  
## [13] 1150000 1250000 9999999 1450000  
## [17] 1500000 6500100 30000000 8900000  
## [21] 300000000 7440000 6782800 2332100  
## [25] 1800000 2000000 10000000 1290000  
## [29] 72456000 34545345 1250000
```

Manipulation application: Summarising data

Summarising data

First we can select only data for Ontario using `filter()`:

```
CES_data %>%  
  filter(cps19_province == "Ontario")
```

```
## # A tibble: 14,160 × 620  
##   cps19_StartDate      cps19_EndDate      cps19_Re...1  
##   <dtm>              <dtm>              <chr>  
## 1 2019-09-13 10:01:19 2019-09-13 10:27:29 R_USWDAPc...  
## 2 2019-09-13 10:05:37 2019-09-13 10:50:53 R_3IQaeDX...  
## 3 2019-09-13 10:05:52 2019-09-13 10:32:53 R_27WeMQ1...  
## 4 2019-09-13 10:10:20 2019-09-13 10:29:45 R_3LiGZcC...  
## 5 2019-09-13 10:14:47 2019-09-13 10:32:32 R_1Iu8R1U...  
## 6 2019-09-13 10:15:39 2019-09-13 10:30:59 R_2EcS26h...  
## 7 2019-09-13 10:15:48 2019-09-13 10:37:45 R_3yrt44w...  
## 8 2019-09-13 10:16:08 2019-09-13 10:40:14 R_100BmXJ...  
## 9 2019-09-13 10:16:24 2019-09-13 10:41:24 R_2e5nvu0...  
## 10 2019-09-13 10:17:06 2019-09-13 10:35:47 R_20Jdv16...  
## # ... with 14,150 more rows, 617 more variables:  
## #   cps19_consent <dbl>, cps19_citizenship <dbl>,  
## #   cps19_yob <dbl>, cps19_yob_2001_age <dbl>,  
## #   cps19_gender <fct>, cps19_province <fct>,  
## #   cps19_education <dbl>, cps19_demsat <dbl>,
```

Summarising data

We don't need to be dealing with all the columns. We can specifically select the ones we want using `select()`:

"How satisfied are you with the performance of your provincial government under \${e://Field/premier}?", "In provincial politics, do you usually think of yourself as a:", and income.

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number)
```

```
## # A tibble: 14,160 × 3  
##   cps19_prov_gov_sat cps19_prov_id cps19...1  
##   <fct>             <fct>         <dbl>  
## 1 Not very satisfied Liberal         NA  
## 2 Fairly satisfied  Progressive Conserva... NA  
## 3 Fairly satisfied  Liberal        56000  
## 4 Not at all satisfied NDP             NA  
## 5 Not at all satisfied NDP              0  
## 6 Not at all satisfied None             NA  
## 7 Not at all satisfied NDP             NA
```

Summarising data

Now that our data looks like what we would like it to, we can start creating a summary table. Since we have the income for each participant, we can look at median incomes. We also want to know how many participants are in each category.

First, we can group the data by provincial political self-ID. To do this, we use `group_by()` to group the data and `summarise()` to produce values for each group we have created. We will start with calculating the `median()` for the incomes. We can add multiple arguments to the `summarise()` argument. `n()` adds a count for each group.

Summarising data

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number) %>%  
  group_by(cps19_prov_gov_sat) %>%  
  summarise(median_income = median(cps19_income_number,  
                                   na.rm = TRUE),  
            count = n())
```

```
## # A tibble: 5 × 3  
##   cps19_prov_gov_sat median_income count  
##   <fct>             <dbl> <int>  
## 1 Very satisfied    80000    872  
## 2 Fairly satisfied  80000   2738  
## 3 Not very satisfied 75000   3212  
## 4 Not at all satisfied 72000   6853  
## 5 Don't know/prefer not to answer 50000    485
```

Grouping

In our table, the satisfaction ratings are ordered alphabetically. We would like them to be ordered logically. We can do this by ordering the factor variable.

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number) %>%  
  mutate(cps19_prov_id = factor(cps19_prov_id,  
                                levels = c("Liberal",  
                                             "Progressive Conservative",  
                                             "NDP",  
                                             "Green",  
                                             "Another party",  
                                             "None",  
                                             "Don't know/prefer not to answer")))
```


Grouping

```
## # A tibble: 14,160 × 3
##   cps19_prov_gov_sat cps19_prov_id cps19...1
##   <fct>             <fct>         <dbl>
## 1 Not very satisfied Liberal          NA
## 2 Fairly satisfied  Progressive Conserva... NA
## 3 Fairly satisfied  Liberal          56000
## 4 Not at all satisfied NDP          NA
## 5 Not at all satisfied NDP            0
## 6 Not at all satisfied None          NA
## 7 Not at all satisfied NDP          NA
## 8 Not very satisfied Liberal         NA
## 9 Not very satisfied NDP          NA
## 10 Not at all satisfied Liberal      NA
## # ... with 14,150 more rows, and abbreviated variable
## #   name 'cps19_income_number'
```

Grouping

And combine this with our table from before:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  mutate(cps19_prov_gov_sat = factor(cps19_prov_gov_sat,
                                     levels = c("Not at all satisfied",
                                                "Not very satisfied",
                                                "Fairly satisfied",
                                                "Very satisfied",
                                                "Don't know/prefer not to answer"),
                                     ordered = TRUE))

group_by(cps19_prov_gov_sat) %>%
  summarise(median_income = median(cps19_income_number,
                                   na.rm = TRUE),
            count = n())
```

```
## # A tibble: 5 × 3
##   cps19_prov_gov_sat median_income count
##   <fct>              <dbl> <int>
## 1 Not at all satisfied    72000  6853
## 2 Not very satisfied     75000  3212
```

Grouping

What happens if we group by political identification instead?

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number) %>%  
  group_by(cps19_prov_id) %>%  
  summarise(median_income = median(cps19_income_number,  
                                   na.rm = TRUE),  
            count = n())
```

```
## # A tibble: 7 × 3  
##   cps19_prov_id median_income count  
##   <fct>          <dbl> <int>  
## 1 Liberal      80000  4607  
## 2 NDP          65000  2413  
## 3 Green        60000   812  
## 4 Progressive Conservative 80000  3629  
## 5 Another party 50000    90  
## 6 None         68000  1367  
## 7 Don't know/prefer not to answer 60000  1242
```

Grouping

We could order the parties in a way that makes more sense:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  mutate(cps19_prov_id = factor(cps19_prov_id,
                                levels = c("Liberal",
                                             "Progressive Conservative",
                                             "NDP",
                                             "Green",
                                             "Another party",
                                             "None",
                                             "Don't know/prefer not to say"),
         )

group_by(cps19_prov_id) %>%
  summarise(median_income = median(cps19_income_number,
                                   na.rm = TRUE),
            count = n())
```

Grouping

```
## # A tibble: 7 × 3
##   cps19_prov_id median_income count
##   <fct>          <dbl> <int>
## 1 Liberal      80000    4607
## 2 Progressive  80000    3629
## 3 NDP          65000    2413
## 4 Green        60000     812
## 5 Another party 50000      90
## 6 None         68000    1367
## 7 Don't know/prefer not to answer 60000    1242
```

Grouping

Or we could sort by median income. We can do that using `arrange()`:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  group_by(cps19_prov_id) %>%
  summarise(median_income = median(cps19_income_number,
                                   na.rm = TRUE),
            count = n()) %>%
  arrange(-median_income)
```

```
## # A tibble: 7 × 3
##   cps19_prov_id median_income count
##   <fct>          <dbl> <int>
## 1 Liberal          80000   4607
## 2 Progressive Conservative 80000   3629
## 3 None             68000   1367
## 4 NDP             65000   2413
## 5 Green           60000    812
## 6 Don't know/prefer not to answer 60000   1242
## 7 Another party    50000    90
```

Grouping

`group_by()` can also have multiple arguments, so we can group by `cps19_prov_gov_sat` and `cps19_prov_id` at the same time:

```
CES_data %>%  
  filter(cps19_province == "Ontario") %>%  
  select(cps19_prov_gov_sat,  
         cps19_prov_id,  
         cps19_income_number) %>%  
  mutate(cps19_prov_id = factor(cps19_prov_id,  
                                levels = c("Liberal",  
                                           "Progressive Conservative",  
                                           "NDP",  
                                           "Green",  
                                           "Another party",  
                                           "None",  
                                           "Don't know/prefer not to answer"),  
          cps19_prov_gov_sat = factor(cps19_prov_gov_sat,  
                                       levels = c("Not at all satisfied",  
                                                  "Not very satisfied",  
                                                  "Fairly satisfied",  
                                                  "Very satisfied",  
                                                  "Don't know/prefer not to answer"))) %>%  
  group_by(cps19_prov_gov_sat, cps19_prov_id) %>%
```

Grouping

This table is less easy to read, though. `spread()` can make a table that is wide rather than long. We specify the **key**, the variable that will become our column names, and the **value**, which will become the values in those columns:

```
CES_data %>%
  filter(cps19_province == "Ontario") %>%
  select(cps19_prov_gov_sat,
         cps19_prov_id,
         cps19_income_number) %>%
  mutate(cps19_prov_id = factor(cps19_prov_id,
                                levels = c("Liberal",
                                             "Progressive Conservative",
                                             "NDP",
                                             "Green",
                                             "Another party",
                                             "None",
                                             "Don't know/prefer not to answer"),
         cps19_prov_gov_sat = factor(cps19_prov_gov_sat,
                                      levels = c("Not at all satisfied",
                                                  "Not very satisfied",
                                                  "Fairly satisfied",
                                                  "Very satisfied",
                                                  "Don't know/prefer not to answer"))
```


Grouping

```
## # A tibble: 7 × 6
##   cps19_prov_id1 Not at all2 Not very3 Fairly4 Very5 Don't know6
##   <fct>          <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 Liberal        80000      80000      79999      79876      60000
## 2 Progressiv...  85000      78000      82000      84000      72000
## 3 NDP            65000      65000      76888      80000      37000
## 4 Green          60000      60000      72750      66000      32500
## 5 Another pa...  40000      48500      73500     150000      52000
## 6 None           62000      74000      69000      66000      43000
## 7 Don't know...  68500      59500      70000      85000      50000
## # ... with abbreviated variable names `cps19_prov_id`,
## #   2`Not at all satisfied`, 3`Not very satisfied`,
## #   4`Fairly satisfied`, 5`Very satisfied`,
## #   6`Don't know/prefer not to answer`
```

Exercises

Exercises

1. Filter the rows in the CES_data dataset where the survey-taker is between 30 and 50 (cps19_age).
2. Filter the rows in the CES_data dataset where the survey-taker answered the cps19_votchoice question (i.e. the cps19_votchoice variable is not NA).
3. Select the variables cps19_age and cps19_province from the CES_data dataset.
4. Select all variables except cps19_province from the CES_data dataset.

Exercises

1. Create a variable in the dataset CES_data that states if a person consumes news content or not (i.e. cps19_news_cons is equal to "0 minutes" or it is not).
2. Modify the variable cps19_income_number in the dataset CES_data so that it is measured in thousands (i.e. divide the income number by 1000).

Exercises

1. Use the CES_data dataset. Group by cps19_votchoice. Find both the median and mean rating of Trudeau (cps19_lead_rating_23):
2. Use the CES_data dataset. Group by cps19_imm and cps19_spend_educ. Find the count for each group.

Exercises

- 1 - Fix this error:

```
CES_data %>%  
  summarise(mean = mean(cps19_age)) %>%  
  group_by(cps19_gender)
```

- 2 - Fix this error:

```
CES_data %>%  
  filter(cps19_vote_choice == "Green Party")
```

Exercises

- 3 - Fix this error:

```
CES_data %>%  
  mutate(cps19_fed_donate = factor(cps19_fed_donate,  
                                   levels = c("Yes",  
                                              "No",  
                                              "Don't know/ Prefer not
```

- 4 - Fix this error:

```
CES_data %>%  
  select(cps19_province  
         cps19_age  
         cps19_gender)
```

Any questions?